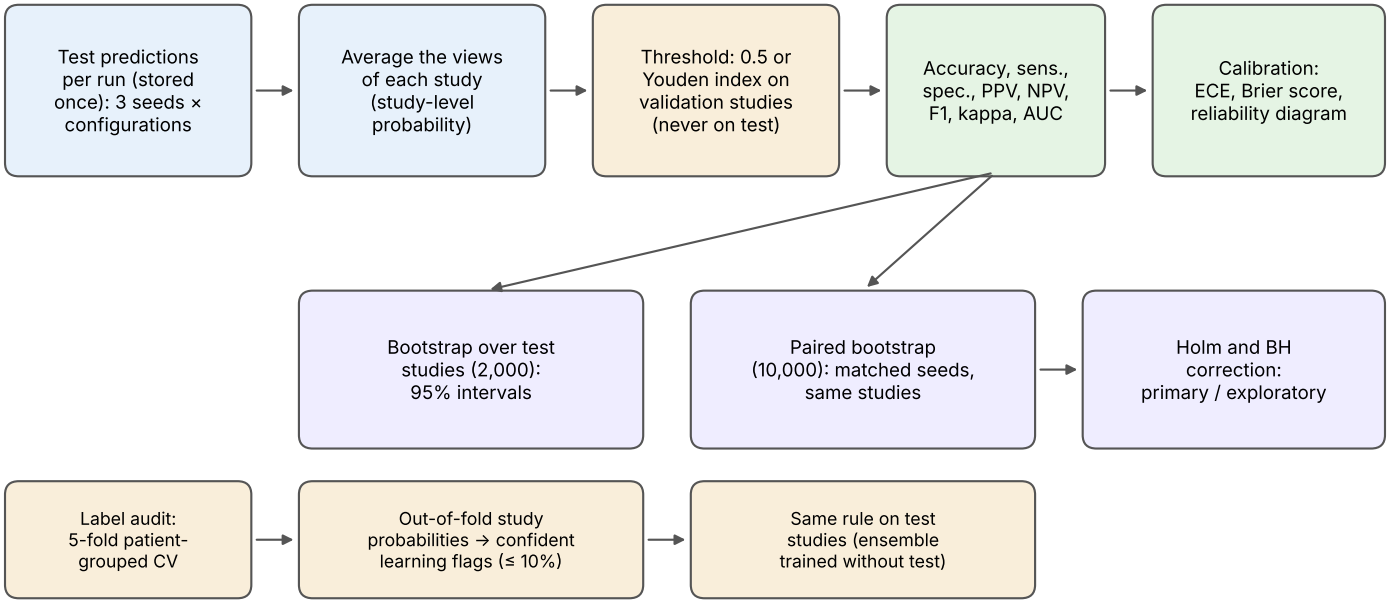
Test predictions per run (stored once): 3 seeds × configurations
Average the views of each study (study-level probability)
Threshold: 0.5 or Youden index on validation studies (never on test)
Accuracy, sens., spec., PPV, NPV, F1, kappa, AUC
Calibration: ECE, Brier score, reliability diagram
Bootstrap over test studies (2,000): 95% intervals
Paired bootstrap (10,000): matched seeds, same studies
Holm and BH correction: primary / exploratory
Label audit: 5-fold patient-grouped CV
Out-of-fold study probabilities → confident learning flags (≤ 10%)
Same rule on test studies (ensemble trained without test)